FOUR-CLASS MVP

Listen, gather context, combine clues

The challenge: useful emotion estimates, soon enough for a conversation.

01 WHAT THE LIVE BASELINE DOES NOW

Solid = current | dashed = next

Capture is continuous. Emotion prediction waits for a completed speech segment.

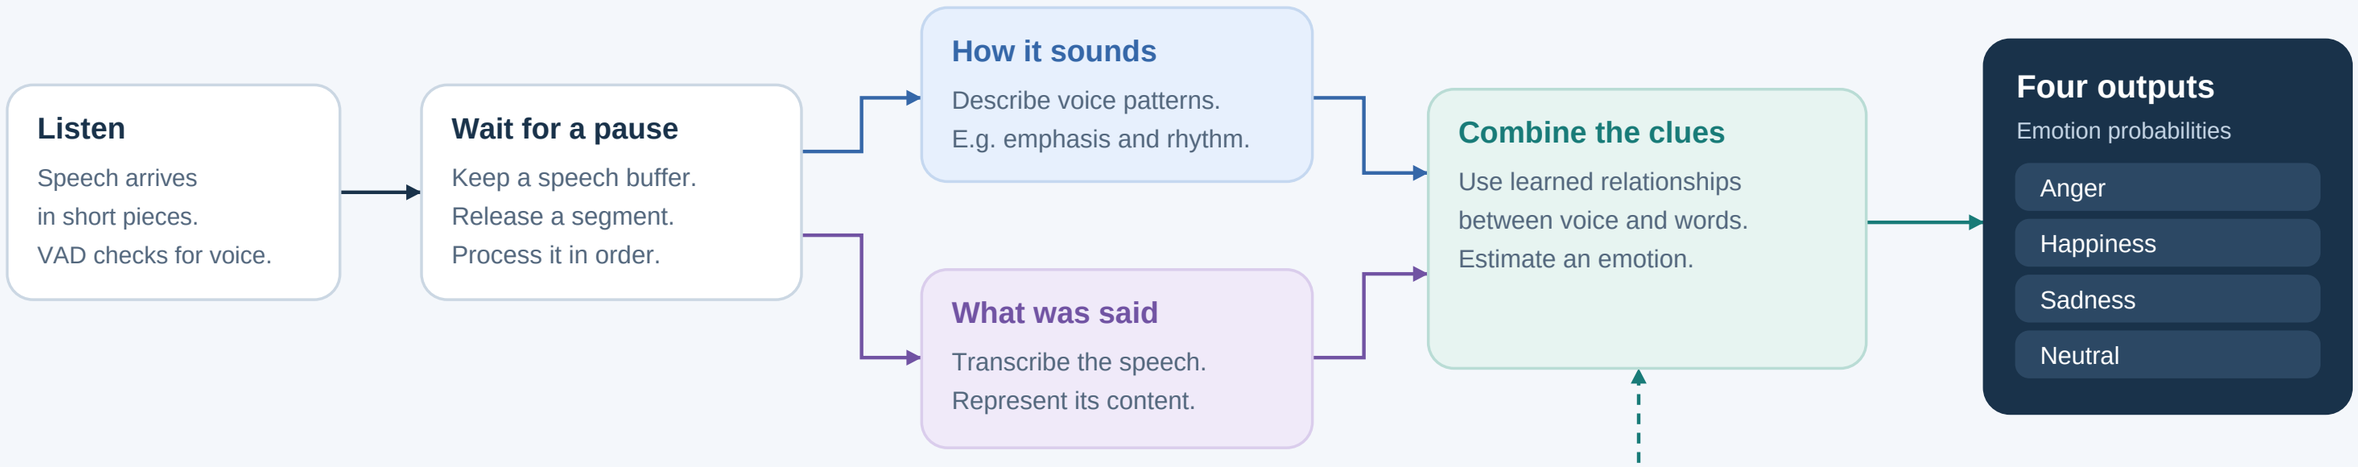

A pause can occur inside a sentence.

Current limit: only the first six seconds of a segment are analyzed.

Long speech needs a deliberate handling policy.

Next: add matching face cues

Features exist; live fusion still needs evaluation.

02 THE DELAY AFTER THE PERSON STOPS SPEAKING

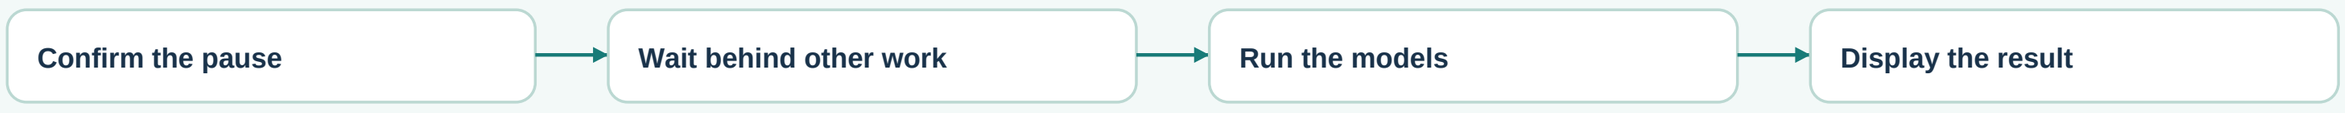

Measure the whole wait. Faster models alone cannot remove pause detection or a backlog.

03 WHEN TO UPDATE IS A DESIGN CHOICE

Rolling windows

Earlier hints; less context.
Overlapping work may repeat.

After a pause (current)

More context; fewer model calls.
Long turns mean more waiting.

Hybrid (proposal)

Tentative updates, then a final result.
More work and coordination.

Better speed must preserve useful emotion predictions. Trust-specific learning is a later task.